

SENTIMENT ANALYSIS FOR

MARKETING

PHASE 4

NAME	SATTANATHAN V
TEAM ID	proj_212173_Team_2

DEVELOPING MODEL PART 2

TOPIC: *Continue building the sentiment analysis solution by Employing NLP techniques, Generating insights.*

Overview of the process:

Sentiment analysis in marketing is a process that involves the use of natural language processing (NLP) techniques to assess and understand the sentiment or emotions expressed in customer feedback, comments, reviews, and other textual data. Here's an overview of the sentiment analysis process for marketing.

Prepare the data: Clean and prepare the data by removing noise, irrelevant information, and special characters.

Tokenize the text into words or phrases for analysis.

Normalize the text (e.g., converting to lowercase) to ensure consistency.

Perform feature selection: This can be done using a variety of methods, such as correlation analysis, information gain, and recursive feature elimination.

Train the model: There are many different machine learning algorithms that can be used for sentiment analysis. Some popular choices include linear regression, random forests, and gradient boosting machines.

1. **Evaluate the model:** This can be done by calculating the mean squared error (MSE) or the root mean squared error (RMSE) of the model's predictions on the held-out test set.
2. **Deploy the model:** Once the model has been evaluated and found to be performing well, it can be deployed to production so that it can be used to predict the house prices of new houses.

PROCEDURE:

Feature selection:

1. **Identify the target variable.** This is the variable that you want to predict, such as house price.
2. **Explore the data.** This will help you to understand the relationships between the different features and the target variable. You can use data visualization and correlation analysis to identify features that are highly correlated with the target variable.
3. **Remove redundant features.** If two features are highly correlated with each other, then you can remove one of the features, as they are likely to contain redundant information.
4. **Remove irrelevant features.** If a feature is not correlated with the target variable, then you can remove it, as it is unlikely to be useful for prediction.

Model training:

1. **Choose a machine learning algorithm.** There are a number of different machine learning algorithms that can be used for sentiment analysis, such as BERT, RoBERTa, Linear regression, random forests, etc.,

Using linear regression:

```
from sklearn.model_selection import train_test_split

from sklearn.linear_model import LogisticRegression

from sklearn.feature_extraction.text import
    TfidfTransformer

from sklearn.feature_extraction.text import
    CountVectorizer

#splitting the data

train_data,test_data,train_labels,test_labels=train_test_split(
    df['text'],df['airline_sentiment'],test_size=0.2,
    random_state=42)

vectorizer = CountVectorizer()

train_data_counts = vectorizer.fit_transform(train_data)

test_data_counts = vectorizer.transform(test_data)
```

```
vec =TfidfTransformer()

train_vec=vec.fit_transform(train_data_counts)

test_vec=vec.transform(test_data_counts)

model=LogisticRegression(max_iter=10000)

model.fit(train_vec,train_labels)
```

OUT:

```
LogisticRegression(max_iter=10000)
```

IN:

```
# Evaluate the model on the test set

accuracy = model.score(test_vec, test_labels)

print(f'Test accuracy: {accuracy:.4f}')
```

OUT:

```
Test accuracy: 0.6282
```

Using Pretrained model BERT

```
import torch

from transformers import
TFBertForSequenceClassification,BertTokenizer,AdamW,
get_linear_schedule_with_warmup,AutoModel,AutoToken
izer,BertModel
```

```
train_data, val_data = train_test_split(df, test_size=0.2,
random_state=42)

tokenizer = BertTokenizer.from_pretrained('bert-base-
uncased')

train_encoding=tokenizer(list(train_data['text']),truncation
=True,padding=True)

valid_encoding=tokenizer(list(val_data['text']),
truncation=True,padding=True)

sentiment_dict = {'positive': 0, 'negative': 1, 'neutral': 2}

train_labels =
train_data['airline_sentiment'].map(sentiment_dict).values.astype
('int64')

valid_labels =
val_data['airline_sentiment'].map(sentiment_dict).values.astype('
int64')

print("train label ",len(train_labels))

print("train label",len(valid_labels))

print("train_encoding ",len(train_encoding))

print("valid_encoding ",len(valid_encoding))

OUT:

train label 11588

train label 2897

train_encoding 3
```

valid_encoding 3

IN:

Split the data into training and validation sets

```
tokenizer = BertTokenizer.from_pretrained('bert-base -  
uncased')
```

Create TensorFlow datasets

```
train_dataset =  
tf.data.Dataset.from_tensor_slices((dict(train_encoding),  
train_labels)).shuffle(len(train_labels)).batch(32)
```

```
val_dataset =  
tf.data.Dataset.from_tensor_slices((dict(valid_encoding),  
valid_labels)).batch(32)
```

Load the pre-trained BERT model for sequence classification

```
model =  
TFBertForSequenceClassification.from_pretrained('bert-base-  
uncased', num_labels=3)
```

Fine-tune the model

```
optimizer = tf.keras.optimizers.Adam(learning_rate=2e-5,  
epsilon=1e-08, clipnorm=1.0)
```

```
loss = tf.keras.losses.SparseCategoricalCrossentropy
```

```
(from_logits=True)
```

```
metric = tf.keras.metrics.SparseCategoricalAccuracy('accuracy')
```

```
model.compile(optimizer=optimizer, loss=loss,  
metrics=[metric])
```

```
history = model.fit(train_dataset, epochs=10,  
validation_data=val_dataset)
```

OUT:

Downloading tf_model.h5: 100%

536M/536M [00:02<00:00, 214MB/s]

All model checkpoint layers were used when initializing
TfBertForSequenceClassification.

Some layers of TfBertForSequenceClassification were not
initialized from the model checkpoint at bert-base-uncased and
are newly initialized: ['classifier']

You should probably TRAIN this model on a down-stream task
to be able to use it for predictions and inference.

Epoch 1/10

363/363 [=====] - 183s
354ms/step - loss: 0.7615 - accuracy: 0.6855 - val_loss: 0.6550 -
val_accuracy: 0.7349

Epoch 2/10

363/363 [=====] - 117s
321ms/step - loss: 0.5929 - accuracy: 0.7590 - val_loss: 0.6105 -
val_accuracy: 0.7432

Epoch 3/10

363/363 [=====] - 116s
319ms/step - loss: 0.4703 - accuracy: 0.8182 - val_loss: 0.6843 -
val_accuracy: 0.7528

Epoch 4/10

363/363 [=====] - 115s
317ms/step - loss: 0.3387 - accuracy: 0.8768 - val_loss: 0.7497 -
val_accuracy: 0.7439

Epoch 5/10

363/363 [=====] - 115s

315ms/step - loss: 0.2372 - accuracy: 0.9152 - val_loss: 0.8090 -
val_accuracy: 0.7539

Epoch 6/10

363/363 [=====] - 117s

322ms/step - loss: 0.1711 - accuracy: 0.9396 - val_loss: 0.9858 -
val_accuracy: 0.7273

Epoch 7/10

363/363 [=====] - 115s

316ms/step - loss: 0.1409 - accuracy: 0.9508 - val_loss: 1.0430 -
val_accuracy: 0.7335

Epoch 8/10

363/363 [=====] - 115s

318ms/step - loss: 0.0993 - accuracy: 0.9669 - val_loss: 1.1365 -
val_accuracy: 0.7366

Epoch 9/10

363/363 [=====] - 114s

315ms/step - loss: 0.0841 - accuracy: 0.9716 - val_loss: 1.2648 -
val_accuracy: 0.7432

Epoch 10/10

363/363 [=====] - 115s

316ms/step - loss: 0.0684 - accuracy: 0.9789 - val_loss: 1.1900 -
val_accuracy: 0.7508

IN:

```
df['airline_sentiment'].value_counts()
```

OUT:

negative 9082

neutral 3069

positive 2334

Name: airline_sentiment, dtype: int64

IN:

```
import seaborn as sns
```

```
sns.countplot(data=df, x='airline_sentiment')
```

OUT:

```
<AxesSubplot:xlabel='airline_sentiment', ylabel='count'>
```


IN:

```
df_new=df.drop(df[df.airline_sentiment  
=='negative'].iloc[:6000].index)
```

```
sns.countplot(data=df_new, x='airline_sentiment')
```

OUT:

```
<AxesSubplot:xlabel='airline_sentiment', ylabel='count'>
```



After Dropping the 6000 data of negative sentiment the datasets seems to be balance

IN:

```
df_new['text']=df_new['text'].apply(preprocess_text)

train_data, test_data = train_test_split(df_new, test_size=0.2,
random_state=42)

tokenizer = BertTokenizer.from_pretrained('bert-base-uncased')

train_encoding=tokenizer(list(train_data['text']),truncation=True,
padding=True)

test_encoding=tokenizer(list(test_data['text']),truncation=True,pa
dding=True)

sentiment_dict = {'positive': 0, 'negative': 1, 'neutral': 2}

train_labels =
train_data['airline_sentiment'].map(sentiment_dict).values.astype
('int64')

test_labels =
test_data['airline_sentiment'].map(sentiment_dict).values.astype(
'int64')

print(len(train_labels))

print(len(test_labels))

print(len(train_encoding))

print(len(test_encoding))
```

OUT:

```
tokenizer = BertTokenizer.from_pretrained('bert-base-uncased')
```

Create TensorFlow datasets

```
train_dataset =  
tf.data.Dataset.from_tensor_slices((dict(train_encoding),  
train_labels)).shuffle(len(train_labels)).batch(32)  
val_dataset =  
tf.data.Dataset.from_tensor_slices((dict(valid_encoding),  
valid_labels)).batch(32)
```

Load the pre-trained BERT model for sequence classification

```
model =  
TFBertForSequenceClassification.from_pretrained('bert-base-  
uncased', num_labels=3)
```

Fine-tune the model

```
optimizer = tf.keras.optimizers.Adam(learning_rate=2e-5,  
epsilon=1e-08, clipnorm=1.0)  
loss =  
tf.keras.losses.SparseCategoricalCrossentropy(from_logits=True  
)  
metric = tf.keras.metrics.SparseCategoricalAccuracy('accuracy')  
model.compile(optimizer=optimizer, loss=loss,  
metrics=[metric])  
history = model.fit(train_dataset, epochs=10,  
validation_data=val_dataset)
```

OUT:

All model checkpoint layers were used when initializing
TFBertForSequenceClassification.

Some layers of TFBertForSequenceClassification were not initialized
from the model checkpoint at bert-base-uncased and are newly
initialized: ['classifier']

You should probably TRAIN this model on a down-stream task to be able to use it for predictions and inference.

Epoch 1/10

213/213 [=====] - 125s 369ms/step -
loss: 0.9499 - accuracy: 0.5514 - val_loss: 0.8267 - val_accuracy:
0.6458

Epoch 2/10

213/213 [=====] - 74s 349ms/step -
loss: 0.7917 - accuracy: 0.6690 - val_loss: 0.7899 - val_accuracy:
0.6631

Epoch 3/10

213/213 [=====] - 72s 336ms/step -
loss: 0.6574 - accuracy: 0.7335 - val_loss: 0.6111 - val_accuracy:
0.7528

Epoch 4/10

213/213 [=====] - 71s 331ms/step -
loss: 0.5206 - accuracy: 0.8005 - val_loss: 0.6245 - val_accuracy:
0.7677

Epoch 5/10

213/213 [=====] - 71s 331ms/step -
loss: 0.3808 - accuracy: 0.8631 - val_loss: 0.6166 - val_accuracy:
0.8001

Epoch 6/10

213/213 [=====] - 71s 332ms/step -
loss: 0.2793 - accuracy: 0.9023 - val_loss: 0.8184 - val_accuracy:
0.7718

Epoch 7/10

213/213 [=====] - 71s 332ms/step -
loss: 0.2128 - accuracy: 0.9262 - val_loss: 0.6901 - val_accuracy:
0.8064

Epoch 8/10

213/213 [=====] - 70s 331ms/step -
loss: 0.1604 - accuracy: 0.9445 - val_loss: 0.9027 - val_accuracy:
0.7829

Epoch 9/10

213/213 [=====] - 71s 332ms/step -
loss: 0.1439 - accuracy: 0.9517 - val_loss: 1.0076 - val_accuracy:
0.7736

Epoch 10/10

213/213 [=====] - 70s 331ms/step -
loss: 0.1116 - accuracy: 0.9613 - val_loss: 0.9026 - val_accuracy:
0.8136

Simple models

from sklearn.naive_bayes **import** MultinomialNB

nb = MultinomialNB()

nb.fit(X_train_tfv,y_train)

MultinomialNB()

from sklearn.linear_model **import** LogisticRegression

log = LogisticRegression(max_iter=1000)

log.fit(X_train_tfv,y_train)

LogisticRegression(max_iter=1000)

from sklearn.svm **import** LinearSVC

svc = LinearSVC()

svc.fit(X_train_tfv,y_train)

LinearSVC()

from sklearn.metrics **import**

plot_confusion_matrix,classification_report

def report(model):

 preds = model.predict(X_test_tfv)

 print(classification_report(y_test,preds))


 plot_confusion_matrix(model,X_test_tfv,y_test)

```
print("NB MODEL")
report(nb)
```

NB MODEL

	precision	recall	f1-score	support
0	0.71	0.98	0.83	1817
1	0.73	0.25	0.37	628
2	0.85	0.36	0.51	483
accuracy		0.72		2928
macro avg	0.77	0.53	0.57	2928
weighted avg	0.74	0.72	0.68	2928

/opt/conda/lib/python3.7/site-packages/sklearn/utils/deprecation.py:87:
FutureWarning: Function plot_confusion_matrix is deprecated; Function
`plot_confusion_matrix` is deprecated in 1.0 and will be removed in 1.2.
Use one of the class methods:
ConfusionMatrixDisplay.from_predictions or
ConfusionMatrixDisplay.from_estimator.
warnings.warn(msg, category=FutureWarning)




```
print("Logistic Regression")
report(log)
```

OUT:

Logistic Regression

	precision	recall	f1-score	support
0	0.79	0.94	0.86	1817
1	0.65	0.44	0.53	628
2	0.83	0.58	0.68	483
accuracy			0.78	2928
macro avg	0.76	0.65	0.69	2928
weighted avg	0.77	0.78	0.76	2928

```
/opt/conda/lib/python3.7/site-
packages/sklearn/utils/deprecation.py:87: FutureWarning:
Function plot_confusion_matrix is deprecated; Function
`plot_confusion_matrix` is deprecated in 1.0 and will be
removed in 1.2. Use one of the class methods:
ConfusionMatrixDisplay.from_predictions or
ConfusionMatrixDisplay.from_estimator.
warnings.warn(msg, category=FutureWarning)
```



```
print('SVC')
report(svc)
```


OUT:

SVC

	precision	recall	f1-score	support
0	0.82	0.89	0.85	1817
1	0.58	0.53	0.55	628
2	0.76	0.62	0.68	483
accuracy		0.77		2928
macro avg	0.72	0.68	0.70	2928
weighted avg	0.76	0.77	0.76	2928

/opt/conda/lib/python3.7/site-packages/sklearn/utils/deprecation.py:87: FutureWarning: Function plot_confusion_matrix is deprecated; Function `plot_confusion_matrix` is deprecated in 1.0 and will be removed in 1.2. Use one of the class methods:

ConfusionMatrixDisplay.from_predictions or
ConfusionMatrixDisplay.from_estimator.
warnings.warn(msg, category=FutureWarning)



Pipeline

```
from sklearn.pipeline import Pipeline
```

```
pipe = Pipeline([('tfidf', TfidfVectorizer()),  
                 ('svc', LinearSVC())])  
pipe.fit(raw_data['text'], raw_data['labels'])
```

```
Pipeline(steps=[('tfidf', TfidfVectorizer()), ('svc', LinearSVC())])
```

```
new_positive_tweet = ['good flight']  
pipe.predict(new_positive_tweet)  
array([2])
```

```
new_negative_tweet = ['bad flight']  
pipe.predict(new_negative_tweet)
```

```
array([0])
```

```
new_neutral_tweet = ['ok flight']  
pipe.predict(new_neutral_tweet)
```

```
array([1])
```

```
##pandasDF --> Hugging Face dataset
```

```
from datasets import Dataset  
dataset = {"text": raw_data["text"].tolist(), "labels":  
raw_data["labels"].tolist()}  
dataset = Dataset.from_dict(dataset)  
dataset = dataset.train_test_split(train_size=0.8, seed=101)  
dataset
```

```
DatasetDict({  
  train: Dataset({  
    features: ['text', 'labels'],  
    num_rows: 11712  
  })  
  test: Dataset({  
    features: ['text', 'labels'],  
    num_rows: 2928  
  })  
})
```

```
IN:
```

```
import tensorflow as tf  
from transformers import  
TFAutoModelForSequenceClassification, AutoTokenizer,  
AutoConfig, DataCollatorWithPadding  
from scipy.special import softmax
```

```
#cardiffnlp/twitter-roberta-base-sentiment
```

```
checkpoint = 'cardiffnlp/twitter-roberta-base-sentiment-latest'
```

```
batch_size = 16
```

```
num_epochs = 5
```

```
config = AutoConfig.from_pretrained(checkpoint)
```

```
tokenizer = AutoTokenizer.from_pretrained(checkpoint)
```

```
model =
```

```
TFAutoModelForSequenceClassification.from_pretrained(checkpoint, num_labels=3)
```

OUT:

```
{"model_id":"c68fd6d69d594f61ba400da869640125","version_major":2,"version_minor":0}
```

```
{"model_id":"0defd486c4dd465ea5ad6a7f166867f0","version_major":2,"version_minor":0}
```

```
{"model_id":"29bda555fd4f4ebc8fc2a88d262c7dce","version_major":2,"version_minor":0}
```

```
{"model_id":"e8f512df42c446eaa5f1e87f84ff18f2","version_major":2,"version_minor":0}
```

```
{"model_id":"e521fb77e7364bf4af34b516f5790788","version_major":2,"version_minor":0}
```

All model checkpoint layers were used when initializing TFRobertaForSequenceClassification.

Some layers of TFRobertaForSequenceClassification were not initialized from the model checkpoint at cardiffnlp/twitter-roberta-base-sentiment-latest and are newly initialized:

```
['classifier']
```

You should probably TRAIN this model on a down-stream task

to be able to use it for predictions and inference.

```
def tokenize_function(example):
```

```
    return tokenizer(example['text'], truncation=True,  
max_length = 35)
```

```
tokenized_datasets = dataset.map(tokenize_function,  
batched=True,)
```

```
data_collator = DataCollatorWithPadding(tokenizer=tokenizer,  
return_tensors="tf")
```

OUT:

```
{"model_id": "01f5e00
```

```
46d6f42f7a15365b6bd246b32", "version_major": 2, "version_min  
or": 0}
```

```
{"model_id": "c3e58a6367704ebb9ed92e43ae966e92", "version_  
major": 2, "version_minor": 0}
```

```
tf_train_dataset = tokenized_datasets["train"].to_tf_dataset(  
    columns=["attention_mask", "input_ids", "token_type_ids"],  
    label_cols=["labels"],  
    shuffle=True,  
    collate_fn=data_collator,  
    batch_size=batch_size  
)
```

```
tf_validation_dataset = tokenized_datasets["test"].to_tf_dataset(  
    columns=["attention_mask", "input_ids", "token_type_ids"],  
    label_cols=["labels"],  
    shuffle=False,  
    collate_fn=data_collator,
```

```
    batch_size=batch_size  
)
```

You're using a `RobertaTokenizerFast` tokenizer. Please note that with a fast tokenizer, using the `__call__` method is faster than using a method to encode the text followed by a call to the `pad` method to get a padded encoding.

```
from tensorflow.keras.optimizers.schedules import  
PolynomialDecay  
from tensorflow.keras.optimizers import Adam
```

```
num_train_steps = len(tf_train_dataset) * num_epochs  
lr_scheduler = PolynomialDecay(  
    initial_learning_rate=5e-5, end_learning_rate=0.0,  
    decay_steps=num_train_steps  
)
```

```
opt = Adam(learning_rate=lr_scheduler)
```

```
loss =  
tf.keras.losses.SparseCategoricalCrossentropy(from_logits=True  
)  
model.compile(optimizer=opt, loss=loss, metrics=['accuracy'])  
model.fit(tf_train_dataset, validation_data=tf_validation_dataset,  
          epochs=3, batch_size=batch_size)
```

Epoch 1/3

```
732/732 [=====] - 145s  
135ms/step - loss: 0.4544 - accuracy: 0.8277 - val_loss: 0.4202 -  
val_accuracy: 0.8361
```

Epoch 2/3

```
732/732 [=====] - 80s  
109ms/step - loss: 0.2816 - accuracy: 0.8984 - val_loss: 0.4345 -
```

val_accuracy: 0.8545

Epoch 3/3

732/732 [=====] - 82s

112ms/step - loss: 0.1579 - accuracy: 0.9468 - val_loss: 0.5486 -

val_accuracy: 0.8446

<keras.callbacks.History at 0x79e6d69a04d0>

IN:

```
from transformers import pipeline
```

```
classifier = pipeline("sentiment-  
analysis",tokenizer=tokenizer,model=model)
```

```
predicted_labels = []
```

```
for text in X_test:
```

```
    result = classifier(text)
```

```
    predicted_label = result[0]['label']
```

```
    predicted_labels.append(predicted_label)
```

```
df = pd.DataFrame(X_test)
```

```
df['predictions'] = predicted_labels
```

```
df['labels'] = df["predictions"].apply(lambda x: 0 if x ==  
"negative" else 1 if x == "neutral" else 2)
```

```
df.head()
```

OUT:

	text	predictions	labels
4814	thanks very excited to see it d	positive	2
150	does that mean you don t have a policy for des...	negative	0
5322	any official word whether flight from bwi to m...	neutral	1
4885	i miss mine terribly a for my th anniversary w...	positive	2
7504	at what time all these passengers were sitting...	negative	0

```

from sklearn.metrics import confusion_matrix
import seaborn as sns
print(classification_report(y_test,df['labels']))

```

```

      precision    recall  f1-score   support

0         0.97      0.98      0.97      1817
1         0.93      0.88      0.91       628
2         0.90      0.95      0.92       483

 accuracy          0.95      2928
 macro avg      0.93      0.94      0.94      2928
weighted avg      0.95      0.95      0.95      2928

```

```


sns.heatmap(confusion_matrix(y_test,df['labels']),cmap='viridis',
annot=True,fmt='d')

```

```

<AxesSubplot:>

```



```
new_positive_tweet = ['good flight']  
classifier(new_positive_tweet)
```

OUT:

```
[{'label': 'positive', 'score': 0.9946355223655701}]
```

```
new_negative_tweet = ['bad flight']  
classifier(new_negative_tweet)
```

OUT:

```
[{'label': 'negative', 'score': 0.9965829253196716}]
```

```
new_neutral_tweet = ['ok flight']  
classifier(new_neutral_tweet)
```

OUT:

```
[{'label': 'neutral', 'score': 0.951370358467102}]
```


CONCLUSION:

Sentiment analysis of Twitter US airline data using the BERT model is a powerful and effective tool for understanding customer opinions and emotions in the airline industry. This approach allows airlines to gain valuable insights into passenger sentiment, which can be pivotal for various aspects of their operations and customer service:

- 1. Improved Customer Service:** By monitoring sentiment, airlines can proactively address customer concerns and issues, leading to better customer experiences.
- 2. Crisis Management:** Sentiment analysis using BERT can help airlines identify and respond to potential PR crises quickly.
- 3. Marketing and Campaigns:** Airlines can fine-tune their marketing strategies based on the sentiments expressed by customers on social media, enabling more targeted and resonant campaigns.
- 4. Product and Service Enhancement:** Understanding customer sentiment provides valuable feedback for improving in-flight services, amenities, and operational aspects.
- 5. Real-time Feedback Loop:** The use of BERT in sentiment analysis ensures that airlines have access to real-time feedback, enabling them to adapt swiftly to customer preferences and concerns.

In essence, sentiment analysis using BERT is a vital tool for

airlines to gauge and react to customer sentiment, thereby enhancing customer satisfaction, refining marketing strategies, and ultimately improving their overall services. It demonstrates the power of NLP and machine learning in gaining insights from vast social media data.